

# Introduction to Distributed File Systems

Harvesting, Storing, and Retrieving Data



How Much Data Do We Make?

# How Much Data?



# JUST HOW BIG IS BIG

---

1.7MB of data is created every second by every person during 2020.

---

In the last two years alone, an astonishing 90% of the world's data has been created.

---

2.5 quintillion bytes of data are produced by humans every day.

---

463 exabytes of data will be generated each day by humans as of 2025.

---

95 million photos and videos are shared every day on Instagram.

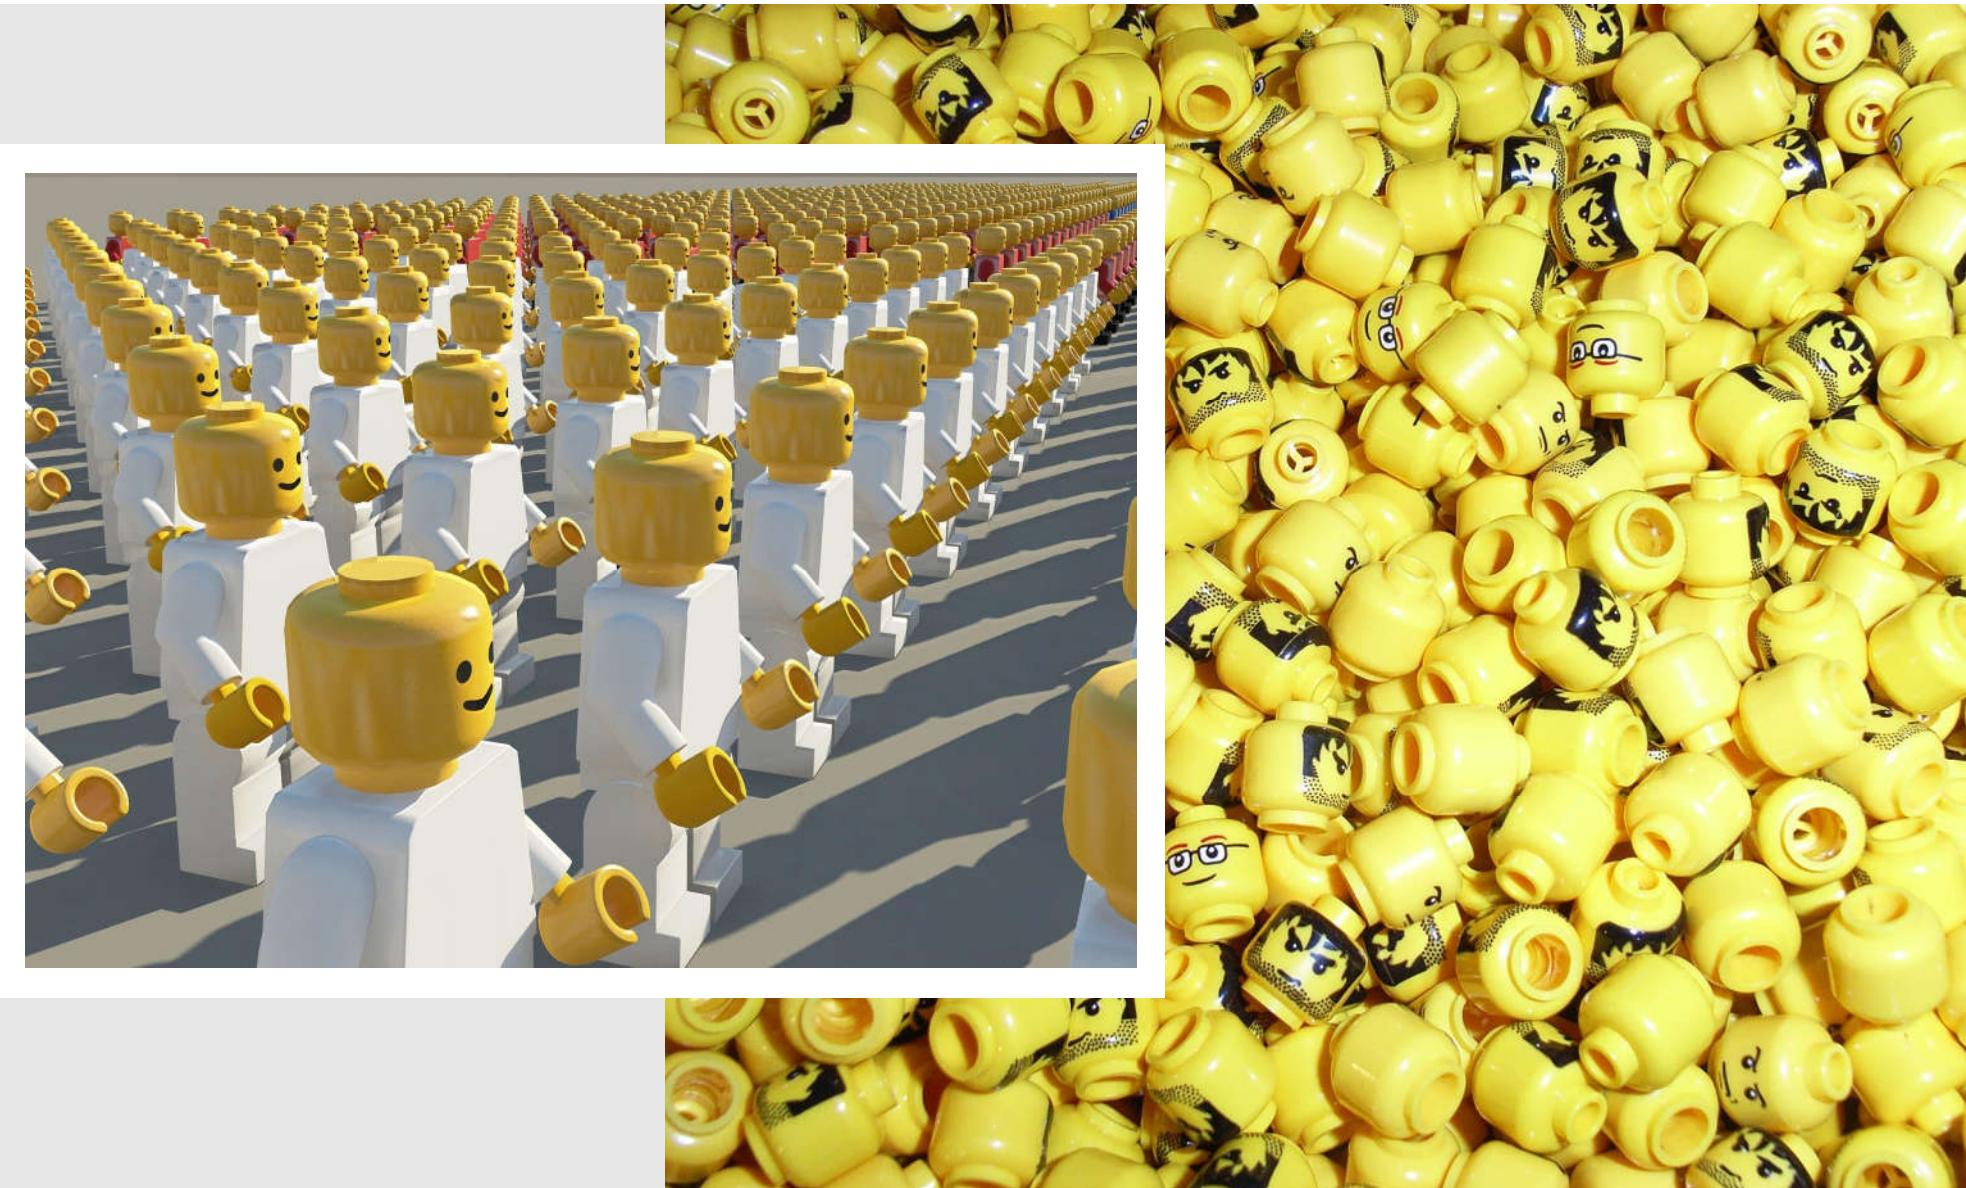
---

By the end of 2020, 44 zettabytes will make up the entire digital universe.

---

Every day, 306.4 billion emails are sent, and 5 million Tweets are made.



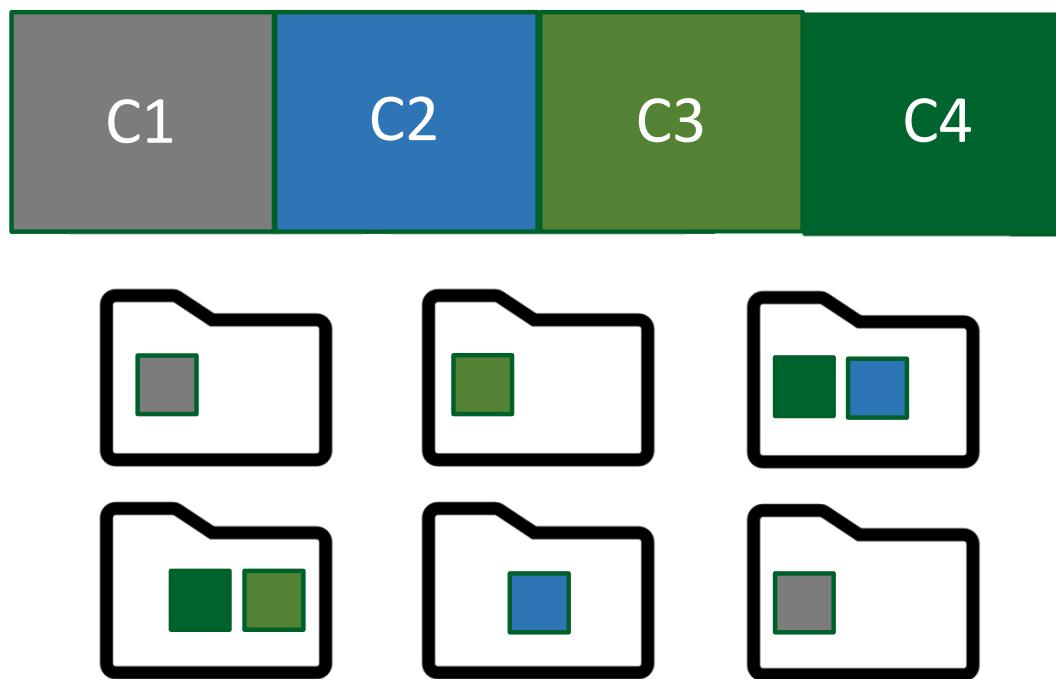




# BIG DATA: HOW TO MANAGE IT?

- Can we use the well-known relational database systems like Oracle, Microsoft, SQL Server, IBM DB2, MySQL, etc., to efficiently manage big data?
- NO!
- Because big data is very different from the data that can be found in relational database systems.
- In other words, big data is NOT the traditional data that is big.
- We need a new type of data management system!

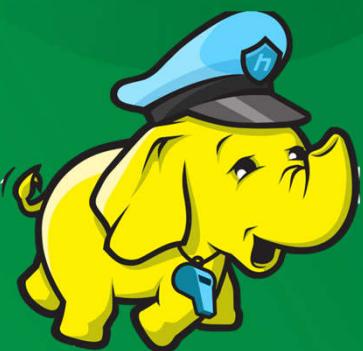
# Google File System



BIG DATA: HOW TO MANAGE IT?  
Enter, Apache Hadoop Framework



# APACHE HADOOP FRAMEWORK: A BIT OF HISTORY



- 2002: Doug Cutting & Mike Cafarella built Nutch to crawl the web. Nutch is a primitive distributed system that could run on several computers at most.
- 2003: Google released the white paper on Google File System that was also a distributed system but much more efficient and powerful.
- 2004: Google released the white paper on MapReduce that is a parallel processing system.
- Based on these white papers, Doug Cutting & Mike Cafarella built a new distributed system that is the foundation for an open-source distributed system: Apache Hadoop Framework.
- And History was made!



# APACHE HADOOP FRAMEWORK: OVERVIEW

## What is Apache Hadoop?

- Apache Hadoop is an open-source framework:
  - Capable of processing large amounts of heterogeneous data sets in a distributed fashion across clusters of commodity computers and hardware
  - Uses a simplified programming model
  - Provides a reliable shared storage and analysis system.



Hadoop Tutorial For Beginners | Hadoop Ecos...

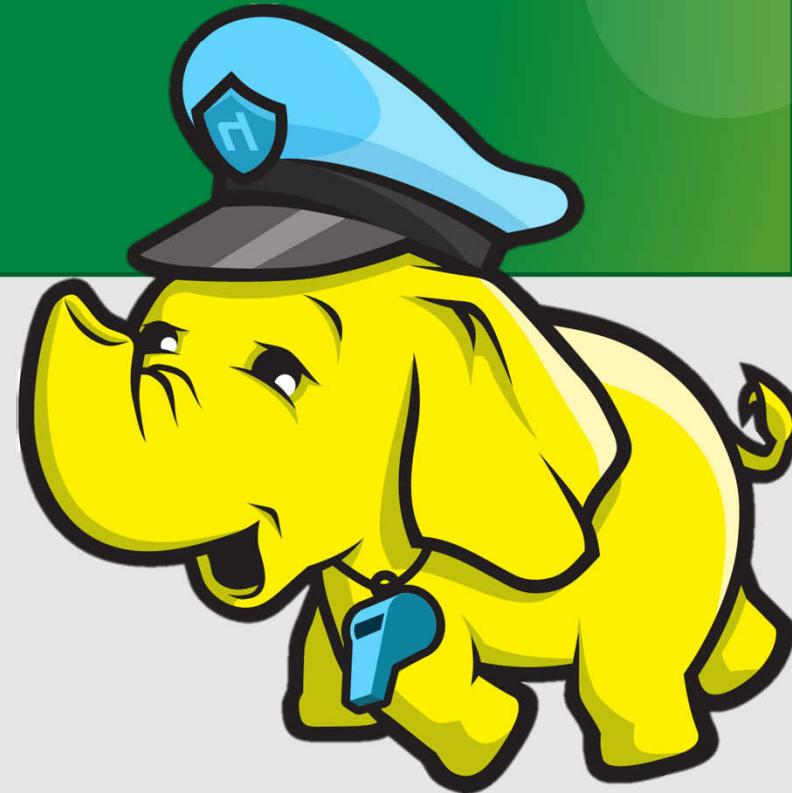
# Hadoop Ecosystem Explained in 20 min!



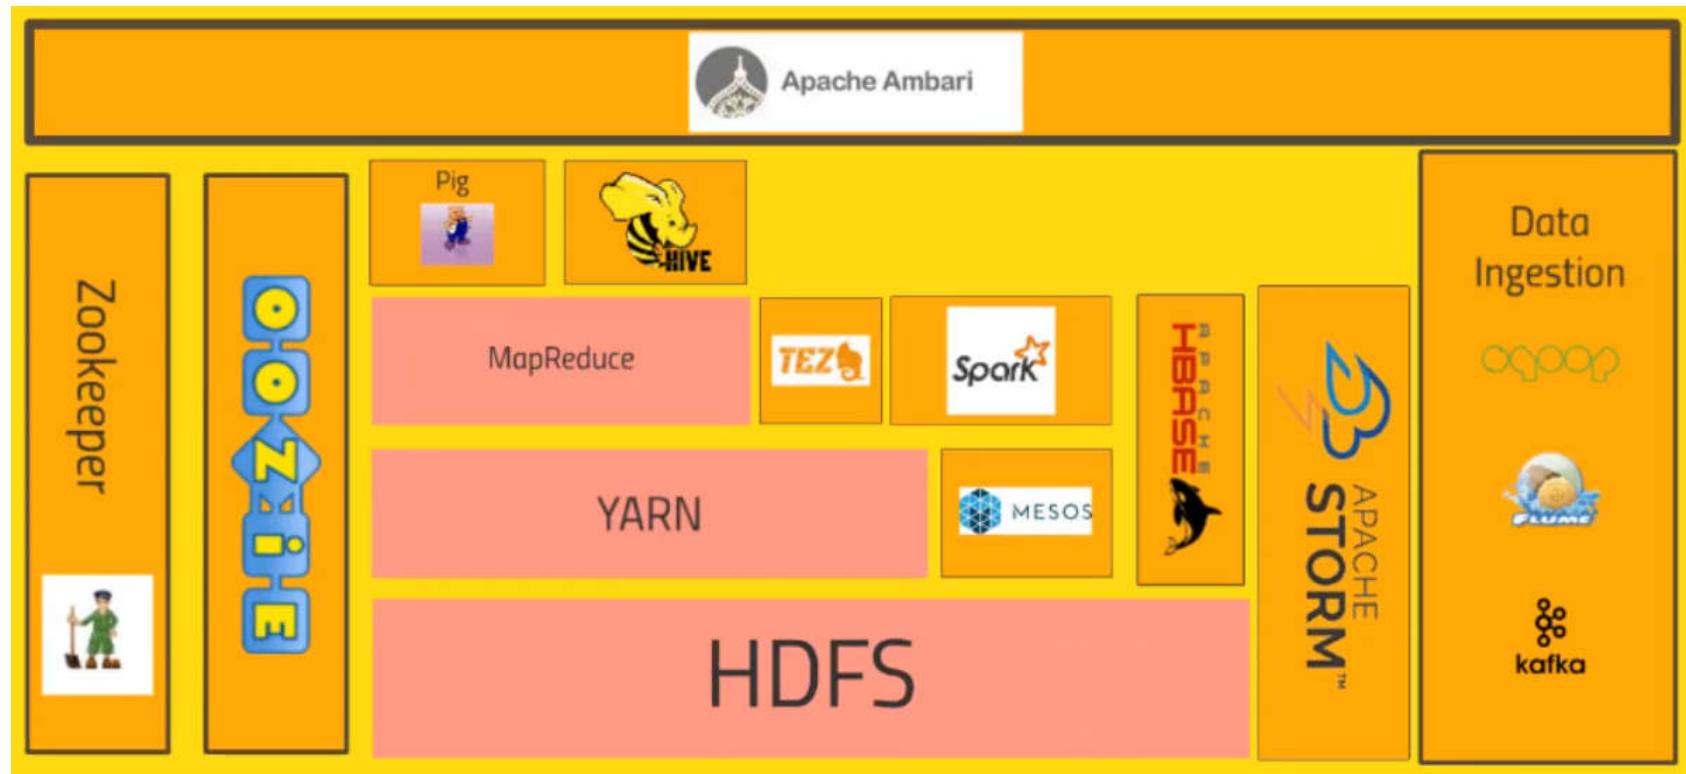
Frank Kane

# APACHE HADOOP ECOSYSTEM: CORE COMPONENTS

- Hadoop Distributed File System (HDFS)
- YARN (Hadoop 2)
- MapReduce
- Hive
- Spark

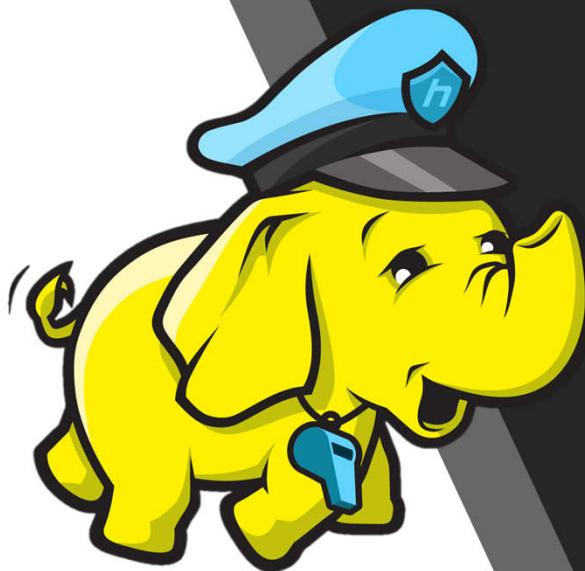


# APACHE HADOOP ECOSYSTEM: CORE COMPONENTS

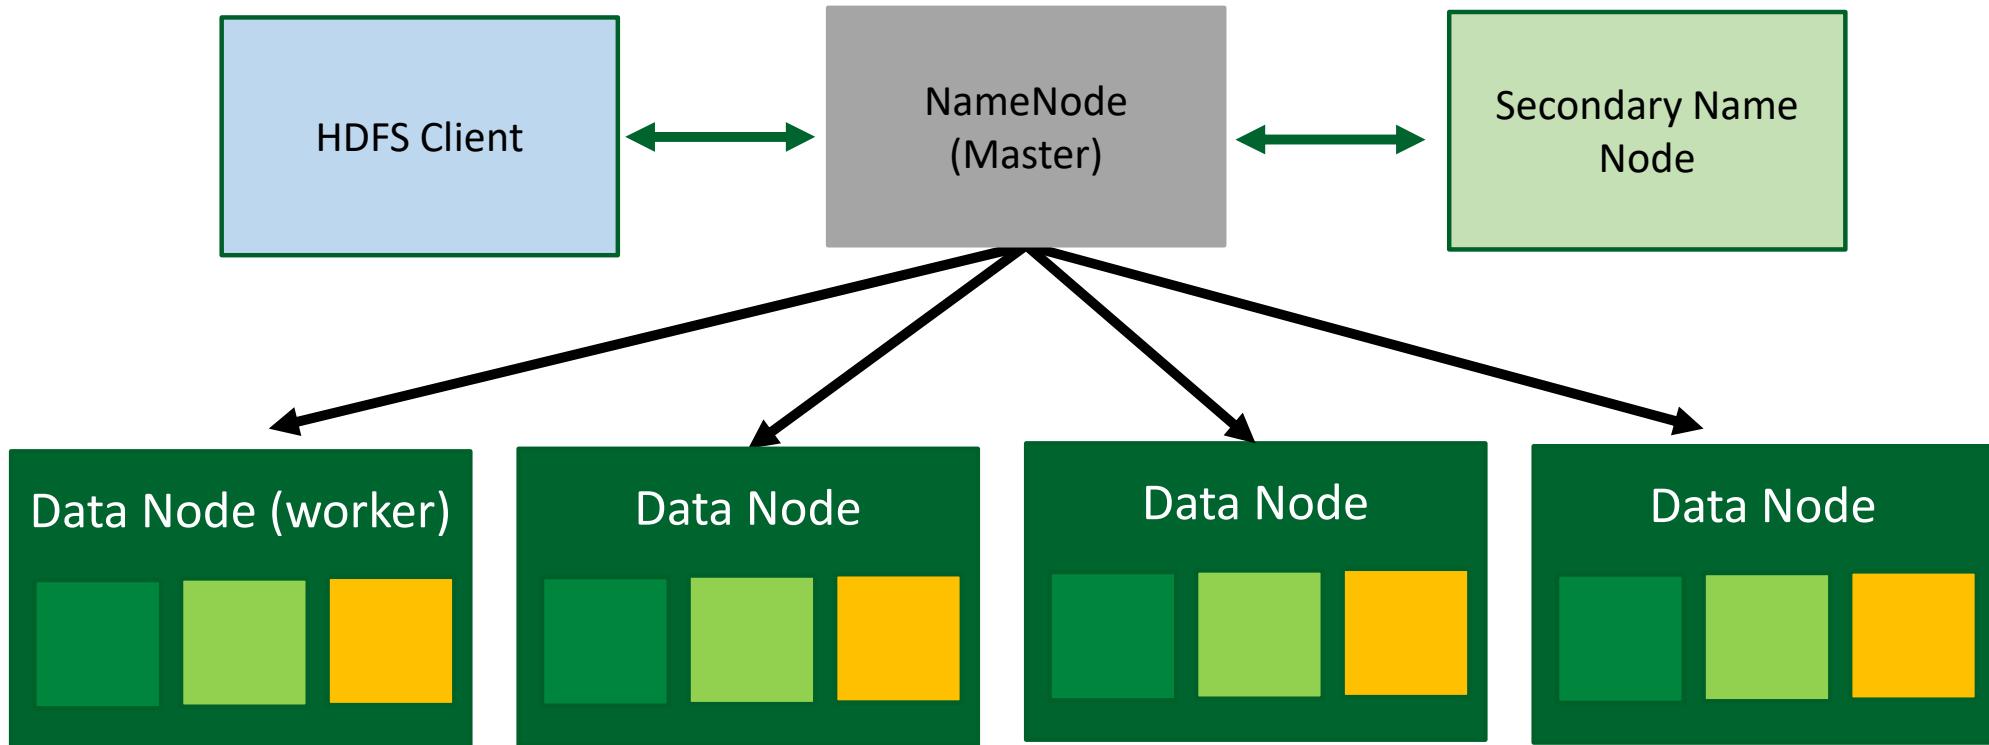


# APACHE HADOOP ECOSYSTEM: HDFS

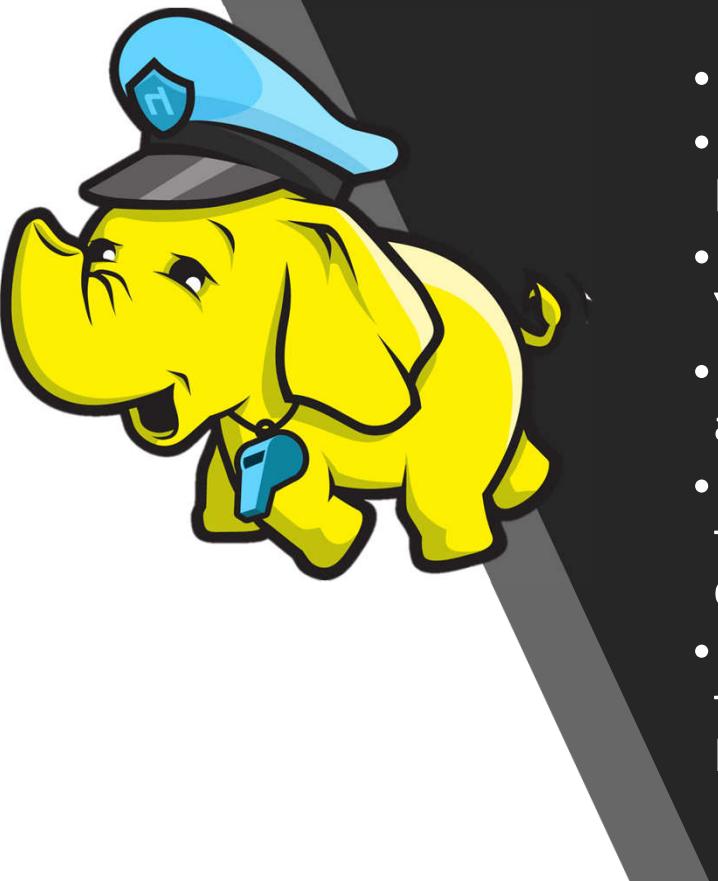
- The "secret sauce"
- Clusters of "computers"
- HDFS is comprised of 3 important components:
  - NameNode
  - DataNode
  - Secondary NameNode
- HDFS operates on a Master-Slave architecture model:
  - NameNode acts as the master node for keeping a track of the storage cluster.
  - DataNode acts as a slave node summing up to the various systems within a Hadoop Cluster



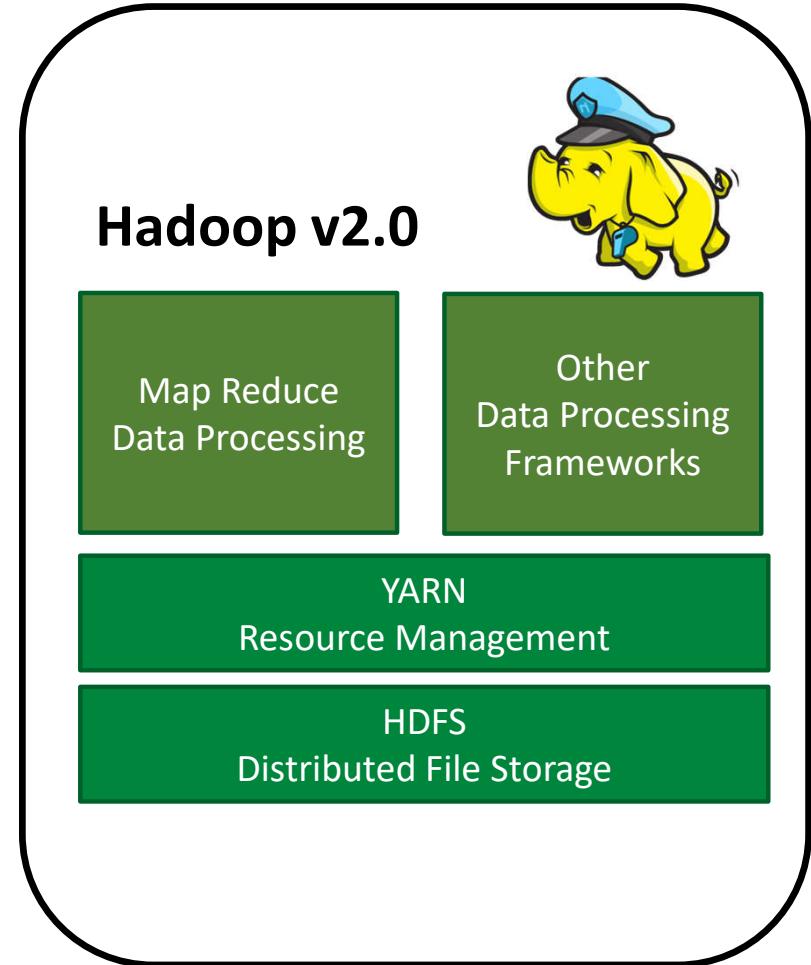
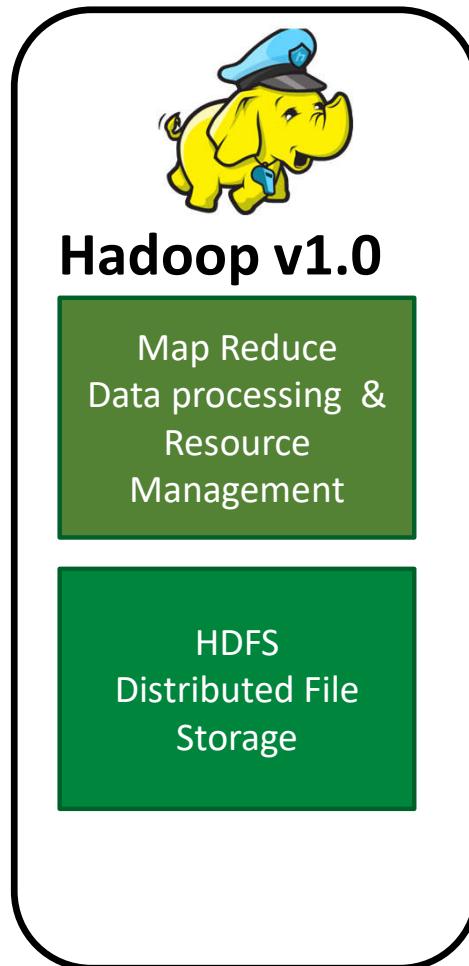
# Apache Hadoop Ecosystem: HDFS Architecture



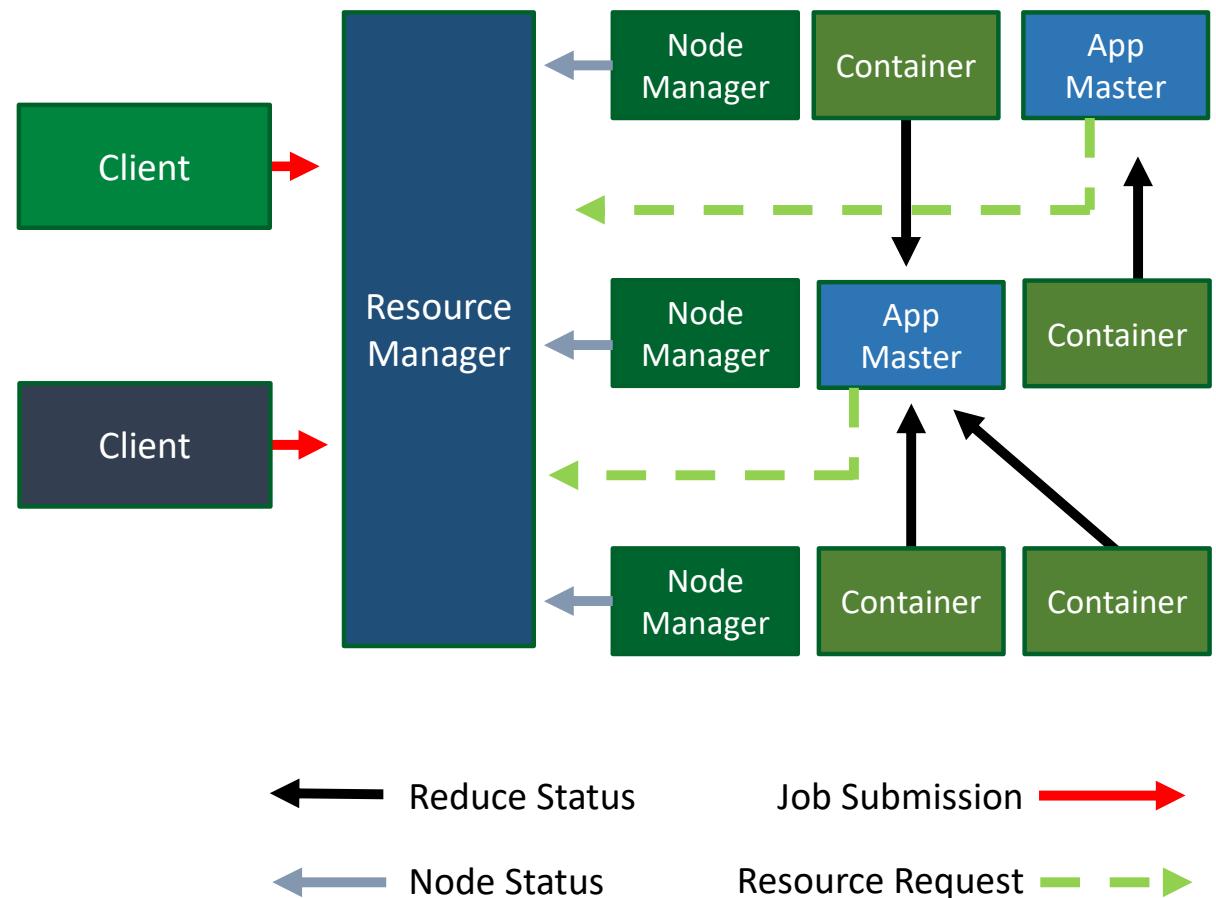
# HADOOP ECOSYSTEM: Yet Another Resource Negotiator (YARN)

- 
- Introduced in Hadoop2
  - Split out from MapReduce, so that alternatives to MapReduce (such as Spark) could be built on top of YARN
  - While HDFS splits up the data storage across your cluster, YARN splits up the computation
  - YARN aligns worker nodes and storage nodes to run jobs as efficiently as possible
  - The resource manager starts a MapReduce application. It first spawns an application master on some node that has capacity.
  - The application master then requests additional nodes that run different application processes (which are usually Map or Reduce tasks).

# APACHE HADOOP ECOSYSTEM: YARN

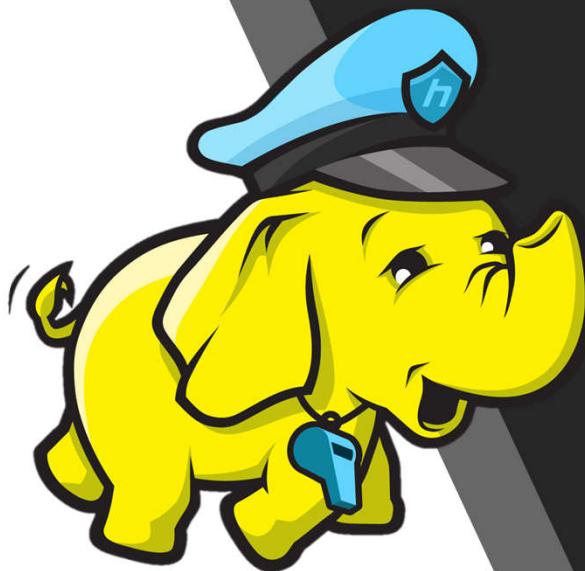


# APACHE HADOOP ECOSYSTEM: YARN Architecture

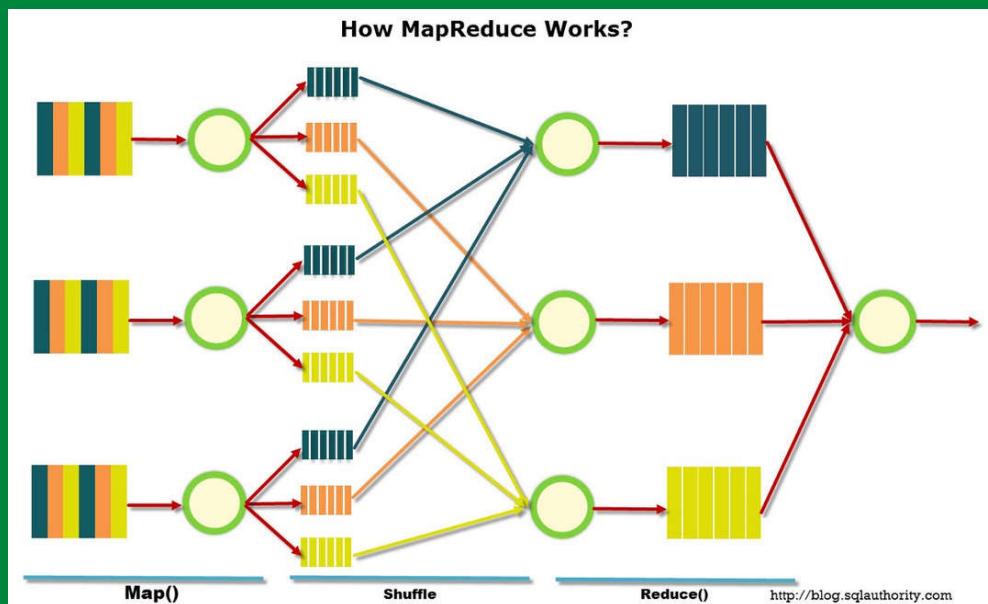


# Apache Hadoop Ecosystem: MapReduce

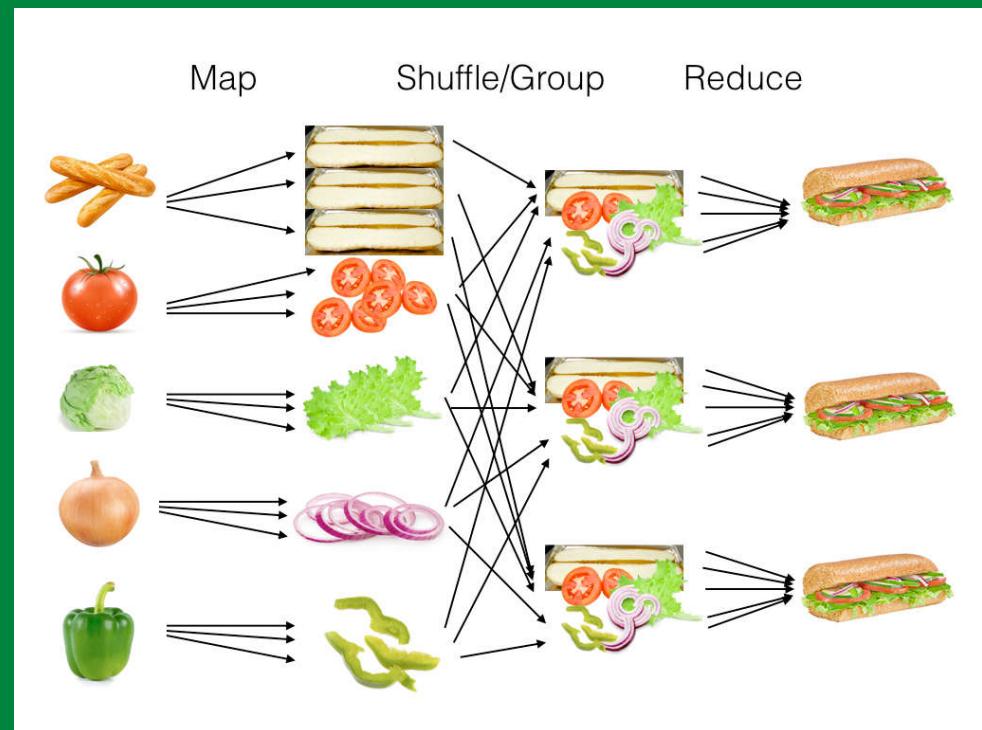
- MapReduce is where data from HDFS gets processed efficiently.
  - Breaks down a big data processing job into smaller tasks
  - Responsible for analyzing large datasets in parallel before finding the results.
- The basic principle of operation behind MapReduce:
  - “Map” job sends a query for processing to various nodes in the cluster.
  - “Reduce” job collects all the results to output into a single value.



# Apache Hadoop Ecosystem: MapReduce

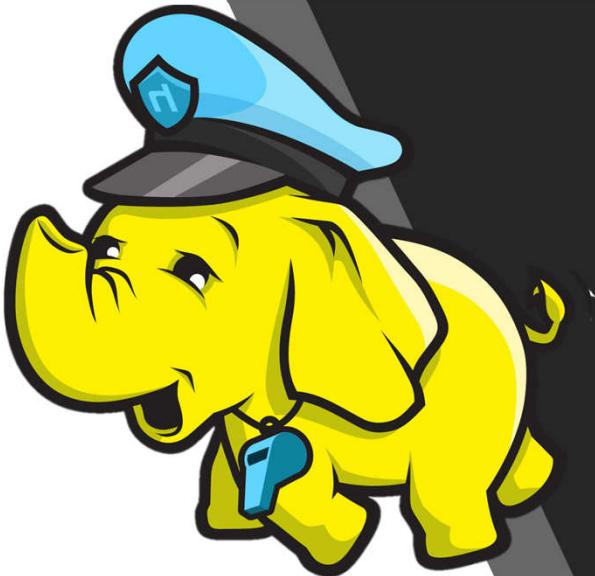


[Image Source](#)



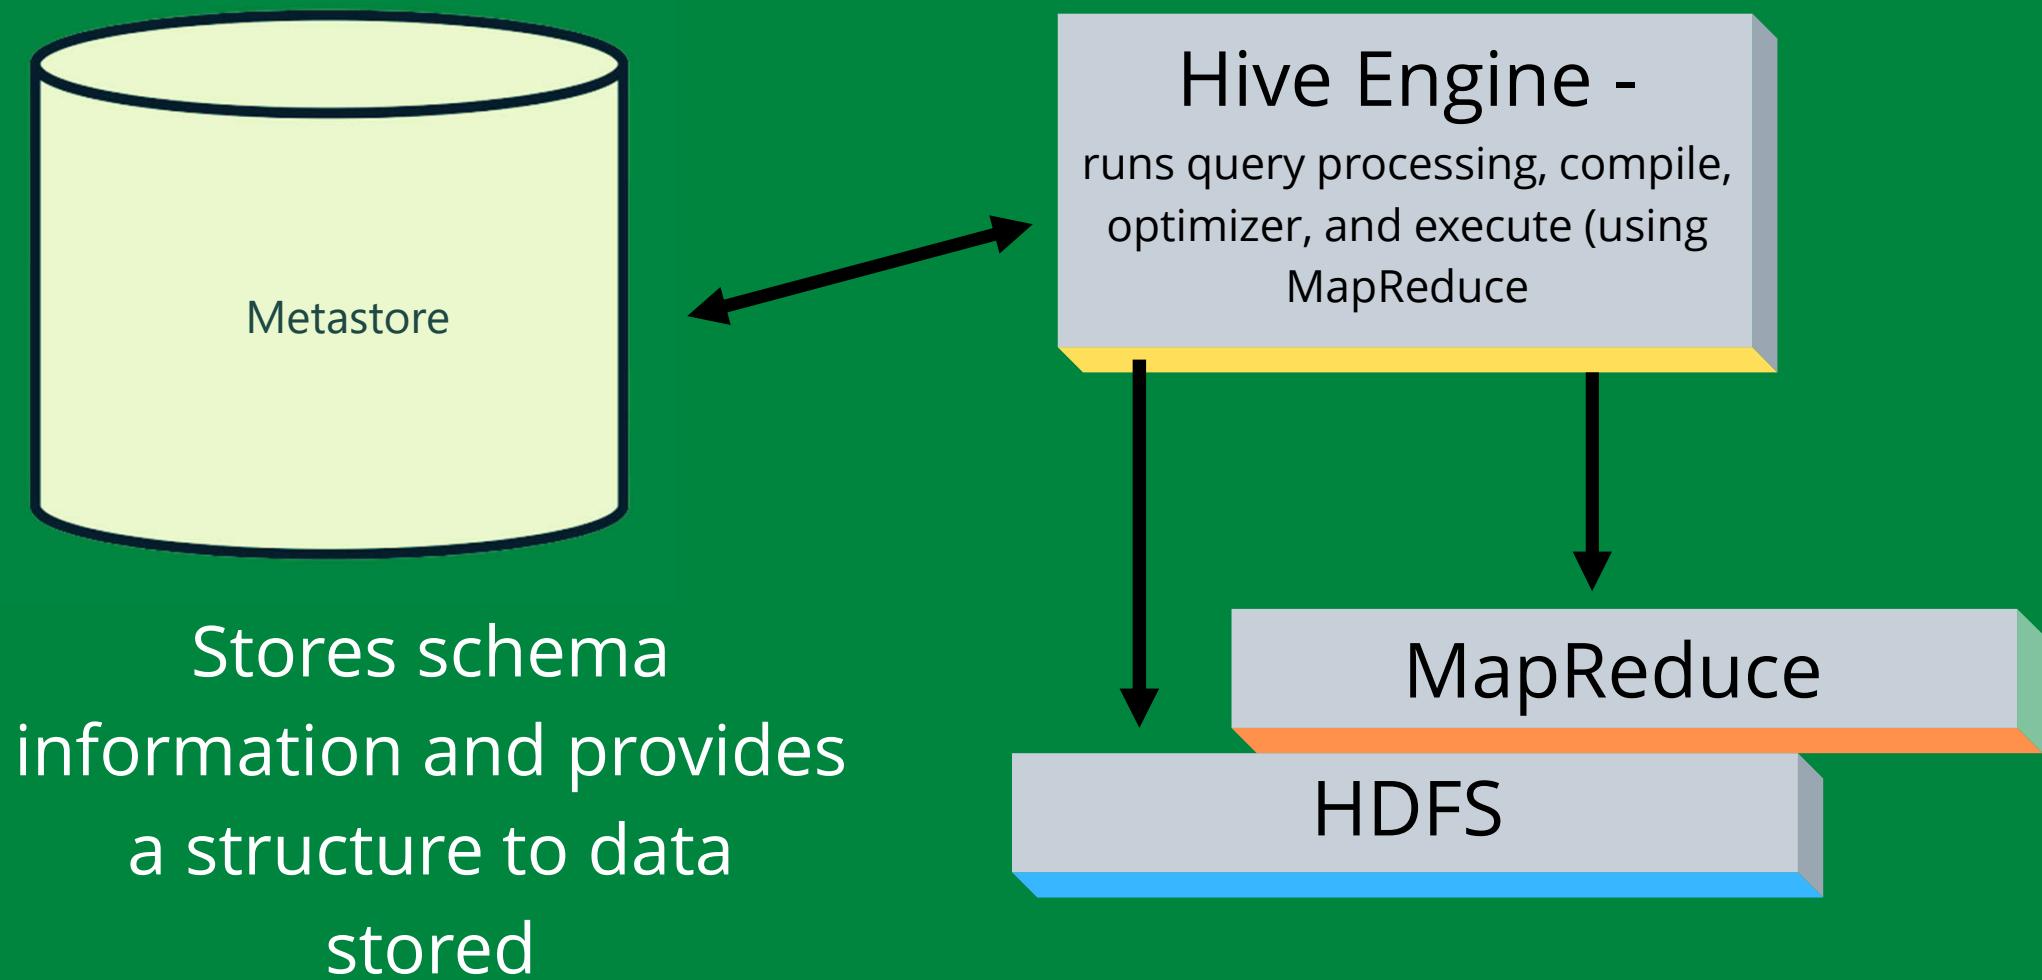
[Image Source](#)

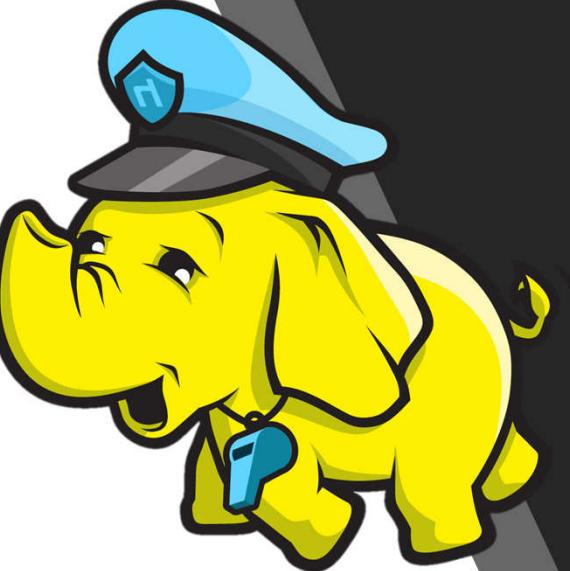
# Apache Hadoop Ecosystem: Hive



- Hive developed by Facebook is a data warehouse:
  - Built on top of Hadoop
  - Provides a simple language known as HiveQL similar to SQL for querying, summarization, and analysis.
  - Hive makes querying faster through indexing.

# Apache Hadoop Ecosystem: Hive - MapReduce - HDFS



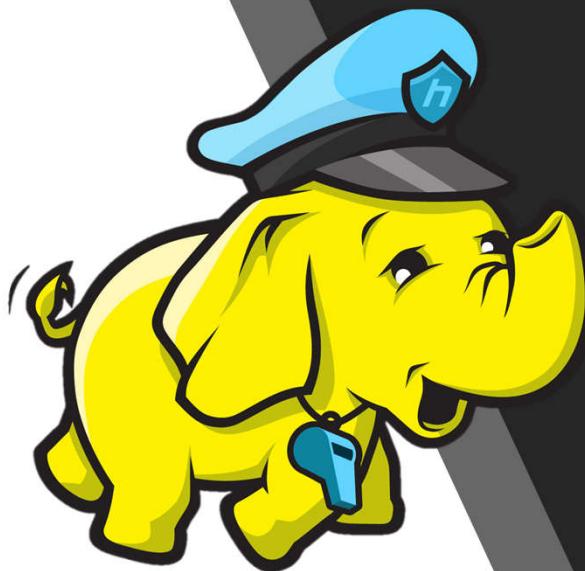


# Apache Hadoop Ecosystem: Spark

- Spark:
  - Like MapReduce: Distribute data across a cluster, and process that data in parallel.
  - Unlike MapReduce which shuffles files around on disk: Spark runs in memory, making much faster at processing data.
- Spark can run on Apache Hadoop clusters, its platforms
- It can access a large variety of data sources:
  - Hadoop Distributed File System
  - (HDFS)
  - Apache
  - Cassandra, Apache HBase
  - Cloud-based storage.

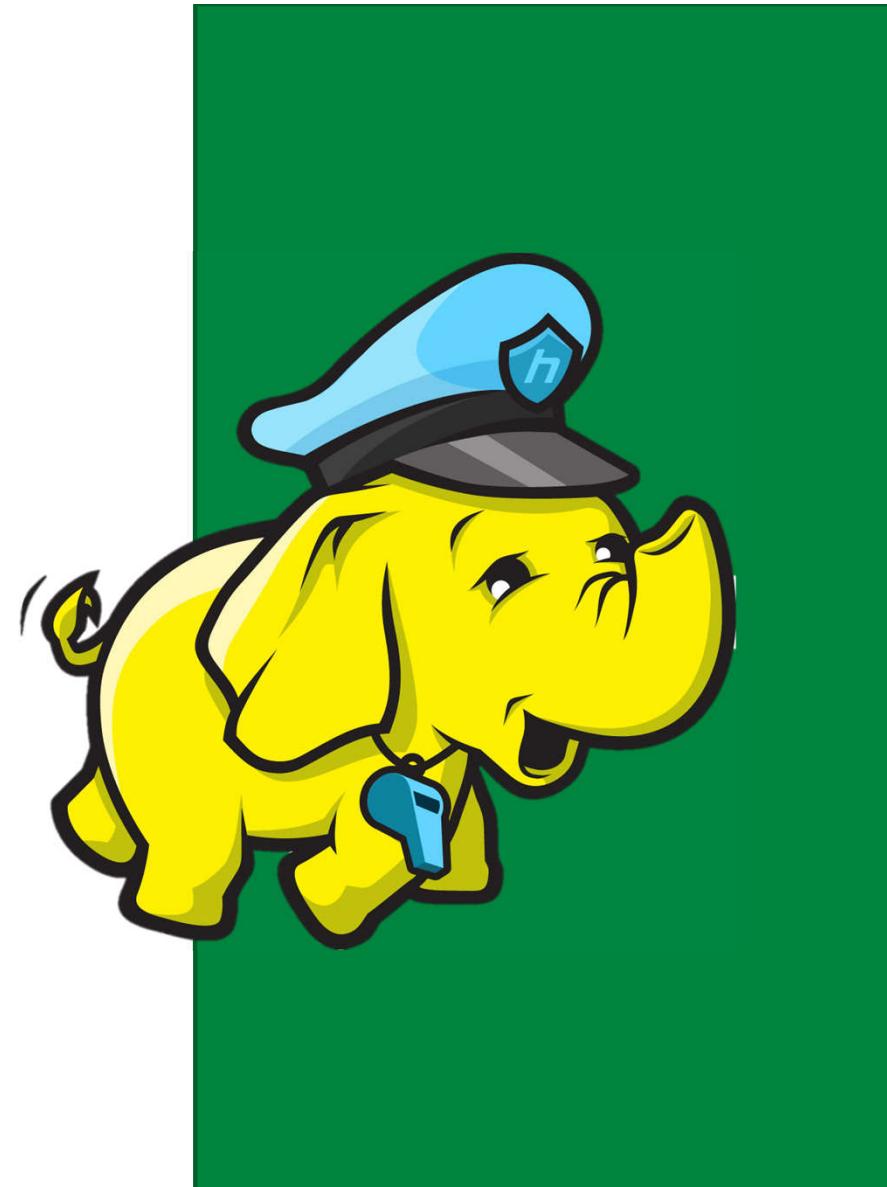
# Apache Hadoop Ecosystem: Spark

- Spark consists of a number of components:
  - Spark Core
  - Spark Streaming
  - Spark Machine Learning Library (MLlib)
  - Spark SQL + DataFrames
  - GraphX

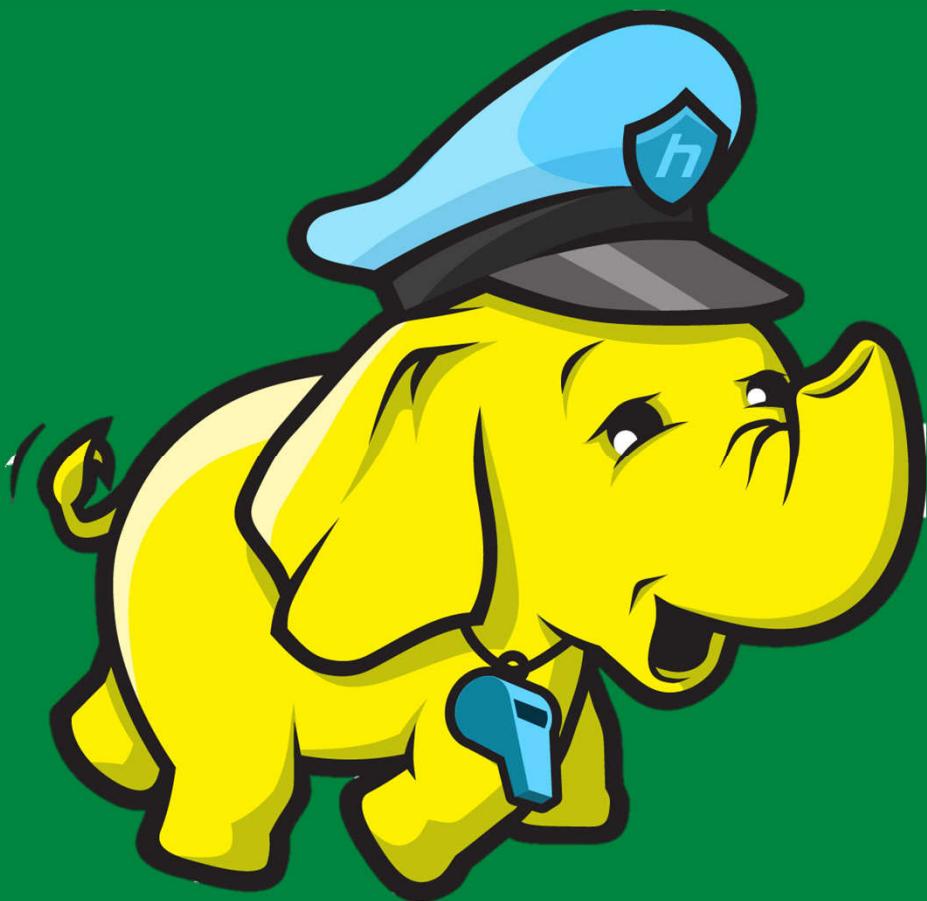


# Apache Hadoop Ecosystem: Why we use it

- Provides reliable shared storage (HDFS) and analysis system (MapReduce/Spark)
- Highly scalable
- Very cost-effective
- Highly flexible
- Has built-in fault tolerance.
- Works on the principle of write once and read multiple times.
- Optimized for large and very large data sets.



## How does Hadoop fit into GCP

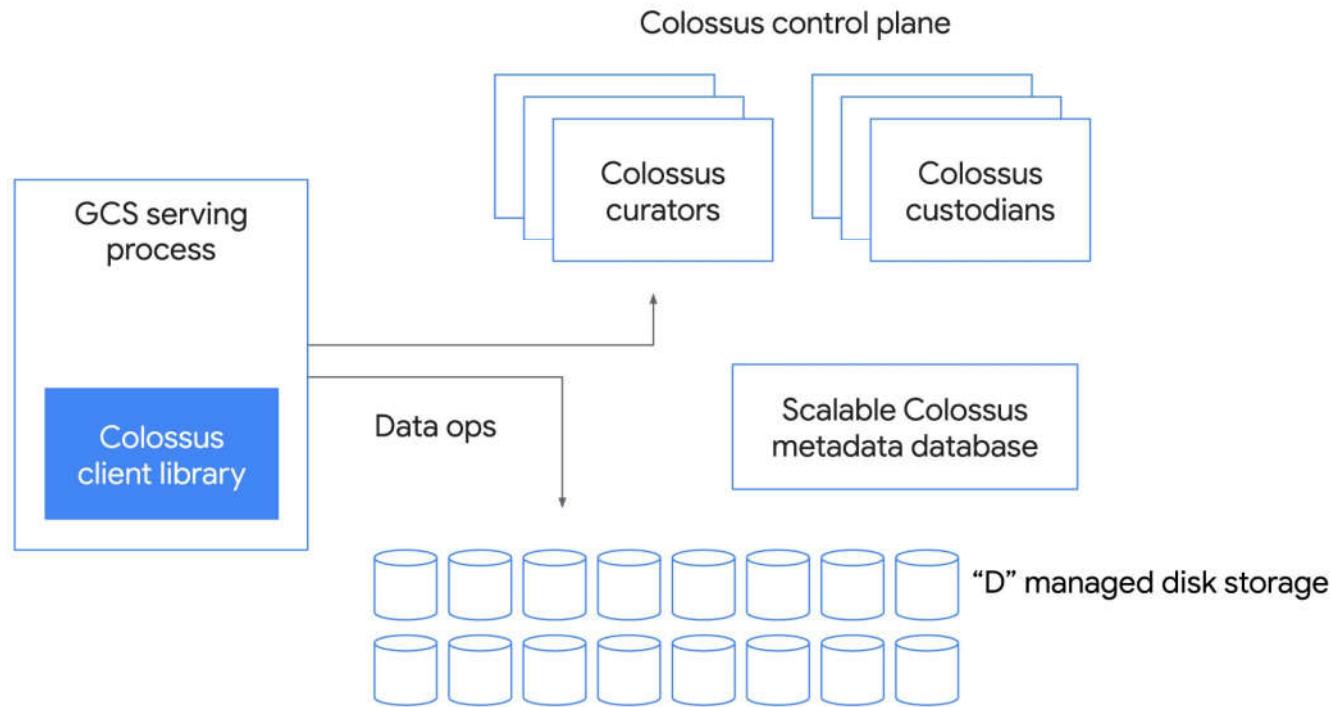




## Understanding Colossus

- Next generation of the Google File System
- Enhances storage scalability
- Improves Availability
- Introduced a distributed metadata model

# Understanding Colossus



# Typical Cluster in Colossus

## Typical cluster

